

```
!pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in  
/usr/local/lib/python3.11/dist-packages (0.13.0)
```

```
Requirement already satisfied: numpy<3,>=1.24.3 in  
/usr/local/lib/python3.11/dist-packages (from imbalanced-learn)  
(2.0.2)
```

```
Requirement already satisfied: scipy<2,>=1.10.1 in  
/usr/local/lib/python3.11/dist-packages (from imbalanced-learn)  
(1.15.3)
```

```
Requirement already satisfied: scikit-learn<2,>=1.3.2 in  
/usr/local/lib/python3.11/dist-packages (from imbalanced-learn)  
(1.6.1)
```

```
Requirement already satisfied: sklearn-compat<1,>=0.1 in  
/usr/local/lib/python3.11/dist-packages (from imbalanced-learn)  
(0.1.3)
```

```
Requirement already satisfied: joblib<2,>=1.1.1 in  
/usr/local/lib/python3.11/dist-packages (from imbalanced-learn)  
(1.5.1)
```

```
Requirement already satisfied: threadpoolctl<4,>=2.0.0 in  
/usr/local/lib/python3.11/dist-packages (from imbalanced-learn)  
(3.6.0)
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from imblearn.over_sampling import SMOTE
```

```
from sklearn.linear_model import LogisticRegression  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.neural_network import MLPClassifier  
from sklearn.metrics import confusion_matrix, classification_report
```

```
df = pd.read_csv('Telco_Customer_Churn.csv')  
df.head()
```

```
{"type": "dataframe", "variable_name": "df"}
```

```
df.drop(['customerID'], axis=1, inplace=True)
```

```
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'],  
errors='coerce')  
df.dropna(inplace=True)
```

```

df['Churn'] = df['Churn'].map({'Yes': 1, 'No': 0})

df = pd.get_dummies(df, drop_first=True)

print("customerID column exists?", 'customerID' in df.columns)

print("TotalCharges type:", df['TotalCharges'].dtype)
print("Missing values in TotalCharges:",
df['TotalCharges'].isnull().sum())

print("Total missing values in dataset:", df.isnull().sum().sum())

print("Unique values in Churn column:", df['Churn'].unique())

print("\nFirst 5 rows of dataframe:")
print(df.head())

```

```

customerID column exists? False
TotalCharges type: float64
Missing values in TotalCharges: 0
Total missing values in dataset: 0
Unique values in Churn column: [0 1]

```

First 5 rows of dataframe:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	Churn
0	0	1	29.85	29.85	0
1	0	34	56.95	1889.50	0
2	0	2	53.85	108.15	1
3	0	45	42.30	1840.75	0
4	0	2	70.70	151.65	1

	Partner_Yes	Dependents_Yes	PhoneService_Yes
0	True	False	False
1	False	False	True
2	False	False	True
3	False	False	False
4	False	False	True

	MultipleLines_No phone service	... StreamingTV_No internet service
0	True	...
1	False	...

False	
2	False ...
False	
3	True ...
False	
4	False ...
False	

	StreamingTV_Yes	StreamingMovies_No internet service
StreamingMovies_Yes \		
0	False	False
False		
1	False	False
False		
2	False	False
False		
3	False	False
False		
4	False	False
False		

	Contract_One year	Contract_Two year	PaperlessBilling_Yes \
0	False	False	True
1	True	False	False
2	False	False	True
3	True	False	False
4	False	False	True

	PaymentMethod_Credit card (automatic) check \	PaymentMethod_Electronic
0	False	
True		
1	False	
False		
2	False	
False		
3	False	
False		
4	False	
True		

	PaymentMethod_Mailed check
0	False
1	True
2	True
3	False
4	False

[5 rows x 31 columns]

```
X = df.drop('Churn', axis=1)
y = df['Churn']

print("Shape of X (features):", X.shape)
print("Shape of y (target):", y.shape)

print("\nFirst 5 rows of X:")
print(X.head())

print("\nFirst 10 values of y:")
print(y.head(10))

print("\nUnique values in y:", y.unique())
```

```
Shape of X (features): (7032, 30)
Shape of y (target): (7032,)
```

First 5 rows of X:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender_Male	\
0	0	1	29.85	29.85	False	
1	0	34	56.95	1889.50	True	
2	0	2	53.85	108.15	True	
3	0	45	42.30	1840.75	True	
4	0	2	70.70	151.65	False	

	Partner_Yes	Dependents_Yes	PhoneService_Yes	\
0	True	False	False	
1	False	False	True	
2	False	False	True	
3	False	False	False	
4	False	False	True	

	MultipleLines_No phone service	MultipleLines_Yes	...	\
0	True	False	...	
1	False	False	...	
2	False	False	...	
3	True	False	...	
4	False	False	...	

	StreamingTV_No internet service	StreamingTV_Yes	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	

	StreamingMovies_No internet service	StreamingMovies_Yes	\
0	False	False	
1	False	False	
2	False	False	

3		False	False
4		False	False

	Contract_One year	Contract_Two year	PaperlessBilling_Yes \
0	False	False	True
1	True	False	False
2	False	False	True
3	True	False	False
4	False	False	True

	PaymentMethod_Credit card (automatic) check \	PaymentMethod_Electronic
0		False
1	True	False
2	False	False
3	False	False
4	False	False

	PaymentMethod_Mailed check
0	False
1	True
2	True
3	False
4	False

[5 rows x 30 columns]

First 10 values of y:

0	0
1	0
2	1
3	0
4	1
5	1
6	0
7	0
8	1
9	0

Name: Churn, dtype: int64

Unique values in y: [0 1]

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
print("Original shape:", X.shape)
print("Scaled shape:", X_scaled.shape)
```

```
Original shape: (7032, 30)
Scaled shape: (7032, 30)
```

```
print("Mean of each column (should be ~0):")
print(np.mean(X_scaled, axis=0))
```


```
print("\nStandard deviation of each column (should be ~1):")
print(np.std(X_scaled, axis=0))
```

```
Mean of each column (should be ~0):
[ 2.62714891e-17 -1.12664271e-16  6.06265133e-17 -1.11906439e-16
 -7.67935835e-17  6.26473971e-17  6.46682808e-17  1.37420097e-16
 -7.27518159e-17  3.63759080e-17  1.97036168e-17 -4.29437802e-18
 -4.29437802e-18 -1.50303231e-17 -4.29437802e-18 -4.40805274e-17
 -4.29437802e-18  7.42674788e-17 -4.29437802e-18  2.72819310e-17
 -4.29437802e-18  8.48771186e-17 -4.29437802e-18  6.66891646e-17
 1.03317683e-16 -1.86931749e-17 -4.06702860e-17 -8.36140662e-17
 5.98686819e-17  4.39542221e-17]
```

[illegible]

```
print("Before SMOTE:")
print(y.value_counts())
sns.countplot(x=y)
plt.title("Class Distribution Before SMOTE")
plt.show()
```

```
Before SMOTE:
Churn
0      5163
1      1869
Name: churn, dtype: int64
```



```
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_scaled, y)

print("After SMOTE:")
print(pd.Series(y_resampled).value_counts())
sns.countplot(x=y_resampled)
plt.title("Class Distribution After SMOTE")
plt.show()
```


After SMOTE:

Churn

0 5163

1 5163

Name: count, dtype: int64



```
X_train, X_test, y_train, y_test = train_test_split(X_resampled,
y_resampled, test_size=0.2, random_state=42)
```

```
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (8260, 30)
X_test shape: (2066, 30)
y_train shape: (8260,)
y_test shape: (2066,)
```

```
lr = LogisticRegression(max_iter=1000)
lr.fit(X_train, y_train)
lr_pred = lr.predict(X_test)
```

```
print("\n Logistic Regression Report:\n")
print(confusion_matrix(y_test, lr_pred))
print(classification_report(y_test, lr_pred))
```

```
\n Logistic Regression Report:
```

```
[[778 259]
```


[[176 853]]					
		precision	recall	f1-score	support
	0	0.82	0.75	0.78	1037
	1	0.77	0.83	0.80	1029
accuracy				0.79	2066
macro avg		0.79	0.79	0.79	2066
weighted avg		0.79	0.79	0.79	2066

```
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)
dt_pred = dt.predict(X_test)
```

```
print("\n Decision Tree Report:\n")
print(confusion_matrix(y_test, dt_pred))
print(classification_report(y_test, dt_pred))
```

\n Decision Tree Report:

[[813 224]					
[221 808]]					
		precision	recall	f1-score	support
	0	0.79	0.78	0.79	1037
	1	0.78	0.79	0.78	1029
accuracy				0.78	2066
macro avg		0.78	0.78	0.78	2066
weighted avg		0.78	0.78	0.78	2066

```
mlp = MLPClassifier(hidden_layer_sizes=(100,), max_iter=300,
random_state=42)
mlp.fit(X_train, y_train)
mlp_pred = mlp.predict(X_test)
```

```
print("\n Neural Network Report:\n")
print(confusion_matrix(y_test, mlp_pred))
print(classification_report(y_test, mlp_pred))
```

\n Neural Network Report:

[[815 222]					
[142 887]]					
		precision	recall	f1-score	support
	0	0.85	0.79	0.82	1037
	1	0.80	0.86	0.83	1029


accuracy			0.82	2066
macro avg	0.83	0.82	0.82	2066
weighted avg	0.83	0.82	0.82	2066



```
/usr/local/lib/python3.11/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic
Optimizer: Maximum iterations (300) reached and the optimization
hasn't converged yet.
warnings.warn(
```

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

def plot_confusion(model_name, y_true, y_pred):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(5, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
    plt.title(f"{model_name} - Confusion Matrix")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

plot_confusion("Logistic Regression", y_test, lr_pred)
plot_confusion("Decision Tree", y_test, dt_pred)
plot_confusion("Neural Network", y_test, mlp_pred)
```





```
from sklearn.metrics import accuracy_score  
models = ['Logistic Regression', 'Decision Tree', 'Neural Network']
```

```


accuracies = [
    accuracy_score(y_test, lr_pred),
    accuracy_score(y_test, dt_pred),
    accuracy_score(y_test, mlp_pred)
]

plt.figure(figsize=(7, 5))
sns.barplot(x=models, y=accuracies, palette="viridis")
plt.title("Model Accuracy Comparison")
plt.ylabel("Accuracy")
plt.ylim(0, 1)
plt.show()

<ipython-input-30-84ae9695587d>:11: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

sns.barplot(x=models, y=accuracies, palette="viridis")

```



```

labels = ['Not Churned', 'Churned']
sizes = y.value_counts()

```

```
plt.figure(figsize=(5, 5))
plt.pie(sizes, labels=labels, autopct='%1.1f%%', startangle=90,
        colors=['#66b3ff', '#ff9999'])
plt.title("Original Churn Distribution")
plt.axis('equal')
plt.show()
```

